

Predicting house prices

August 30, 2017

In []: Aluno: Marco Antonio Santo

1 Fire up libraries

```
In [1]: import matplotlib
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import linear_model
```

2 Load some house sales data

Dataset is from house sales in King County, the region where the city of Seattle, WA is located.

```
In [2]: sales = pd.read_csv('home_data.csv')
```

```
In [3]: sales.head()
```

```
Out[3]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900	3	1.00	1180	
1	6414100192	20141209T000000	538000	3	2.25	2570	
2	5631500400	20150225T000000	180000	2	1.00	770	
3	2487200875	20141209T000000	604000	4	3.00	1960	
4	1954400510	20150218T000000	510000	3	2.00	1680	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	0	0	...	7	1180	
1	7242	2.0	0	0	...	7	2170	
2	10000	1.0	0	0	...	6	770	
3	5000	1.0	0	0	...	7	1050	
4	8080	1.0	0	0	...	8	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0	1955	0	98178	47.5112	-122.257	
1	400	1951	1991	98125	47.7210	-122.319	

2	0	1933	0	98028	47.7379	-122.233
3	910	1965	0	98136	47.5208	-122.393
4	0	1987	0	98074	47.6168	-122.045

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 21 columns]

```
In [4]: sales[sales['id']==1839920160]
```

```
Out[4]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
11860	1839920160	20140714T000000	432000	3	2.0	1870	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
11860	7080	1.0	0	0	...	7	1210	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
11860	660	1969	0	98034	47.7244	-122.179	

	sqft_living15	sqft_lot15
11860	1620	8000

[1 rows x 21 columns]

```
In [5]: sales.keys()
```

```
Out[5]: Index([u'id', u'date', u'price', u'bedrooms', u'bathrooms', u'sqft_living',
              u'sqft_lot', u'floors', u'waterfront', u'view', u'condition', u'grade',
              u'sqft_above', u'sqft_basement', u'yr_built', u'yr_renovated',
              u'zipcode', u'lat', u'long', u'sqft_living15', u'sqft_lot15'],
              dtype='object')
```


```
In [6]: sales.shape
```

```
Out[6]: (21613, 21)
```

3 Exploring the data for housing sales

The house price is correlated with the number of square feet of living space.

```
In [7]: plt.grid('on')
        plt.scatter(sales['sqft_living'], sales['price'])
        plt.show()
```



4 Create a simple regression model of sqft_living to price

Split data into training and testing.

We use `random_state=200` so that everyone running this notebook gets the same results. In practice, you may set a random seed.

```
In [8]: train_data = sales.sample(frac=0.8, random_state=200)
        test_data  = sales.drop(train_data.index)
        print(train_data.shape, test_data.shape)
```

```
((17290, 21), (4323, 21))
```

```
In [9]: train_data.head()
```

```
Out[9]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
11860	1839920160	20140714T000000	432000	3	2.00	1870	
12446	6705850140	20141009T000000	750000	4	2.75	3170	
10556	924069190	20140819T000000	440000	3	1.75	2000	
4828	3211270170	20140523T000000	404000	4	3.00	4060	
3502	9523103001	20141013T000000	389000	2	1.00	850	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
11860	7080	1.0	0	0	...	7	1210	
12446	7634	2.0	0	0	...	10	3170	

10556	11880	2.0	0	0	...	8	2000
4828	35621	1.0	0	0	...	9	2030
3502	3276	1.0	0	0	...	6	850

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
11860	660	1969	0	98034	47.7244	-122.179	
12446	0	1992	0	98075	47.5774	-122.054	
10556	0	1979	0	98075	47.5882	-122.052	
4828	2030	1989	0	98092	47.3059	-122.108	
3502	0	1910	0	98103	47.6742	-122.350	

	sqft_living15	sqft_lot15
11860	1620	8000
12446	2940	7846
10556	1820	15120
4828	2950	35259
3502	1460	4100

[5 rows x 21 columns]

In [10]: test_data.head()

Out[10]:

	id	date	price	bedrooms	bathrooms	sqft_living	\
3	2487200875	20141209T000000	604000	4	3.0	1960	
4	1954400510	20150218T000000	510000	3	2.0	1680	
5	7237550310	20140512T000000	1225000	4	4.5	5420	
17	6865200140	20140529T000000	485000	4	1.0	1600	
18	16000397	20141205T000000	189000	2	1.0	1200	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
3	5000	1.0	0	0	...	7	1050	
4	8080	1.0	0	0	...	8	1680	
5	101930	1.0	0	0	...	11	3890	
17	4300	1.5	0	0	...	7	1600	
18	9850	1.0	0	0	...	7	1200	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
3	910	1965	0	98136	47.5208	-122.393	
4	0	1987	0	98074	47.6168	-122.045	
5	1530	2001	0	98053	47.6561	-122.005	
17	0	1916	0	98103	47.6648	-122.343	
18	0	1921	0	98002	47.3089	-122.210	

	sqft_living15	sqft_lot15
3	1360	5000
4	1800	7503
5	4760	101930
17	1610	4300

18 1060 5095

[5 rows x 21 columns]


4.1 Build the regression model using only sqft_living as a feature

```
In [11]: x_train = train_data['sqft_living'].values.reshape(-1,1)
         y_train = train_data['price'].values.reshape(-1,1)
```

```
In [12]: simple_model = linear_model.LinearRegression()
         simple_model.fit(x_train, y_train)
```

```
Out[12]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [13]: plt.grid('on')
         plt.scatter(x_train, y_train)
         plt.plot(x_train, simple_model.predict(x_train), color='red', linewidth=2)
         plt.show()
```




5 Let's show what our predictions look like

```
In [14]: x_test = test_data['sqft_living'].values.reshape(-1,1)
         y_test = test_data['price'].values.reshape(-1,1)

         y_pred = simple_model.predict(x_test)
```

```
In [15]: plt.grid('on')
plt.scatter(x_test, y_test)
plt.plot(x_test, y_pred, color='red', linewidth=2)
plt.show()
```



6 Evaluate the simple model

```
In [16]: def rmse(predictions, targets):
return np.sqrt(((predictions - targets) ** 2).mean())

In [17]: print('intercept:', simple_model.intercept_, 'coefficients:', simple_model.coef_)
# The mean squared error
print("RMSE: %.2f" % (rmse(y_pred, y_test)))

('intercept:', array([-46493.04519733]), 'coefficients:', array([[ 282.27187583]]))
RMSE: 254323.39
```

RMSE of about \$254,323,39

7 Explore other features in the data

To build a more elaborate model, we will explore using more features.

```
In [18]: my_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode']
```


```
In [19]: sales[my_features].describe()
```

```
Out [19]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors \
count	21613.000000	21613.000000	21613.000000	2.161300e+04	21613.000000
mean	3.370842	2.114757	2079.899736	1.510697e+04	1.494309
std	0.930062	0.770163	918.440897	4.142051e+04	0.539989
min	0.000000	0.000000	290.000000	5.200000e+02	1.000000
25%	3.000000	1.750000	1427.000000	5.040000e+03	1.000000
50%	3.000000	2.250000	1910.000000	7.618000e+03	1.500000
75%	4.000000	2.500000	2550.000000	1.068800e+04	2.000000
max	33.000000	8.000000	13540.000000	1.651359e+06	3.500000

	zipcode
count	21613.000000
mean	98077.939805
std	53.505026
min	98001.000000
25%	98033.000000
50%	98065.000000
75%	98118.000000
max	98199.000000

```
In [20]: #sales.show(view='BoxWhisker Plot', x='zipcode', y='price')
plt.grid('on')
plt.scatter(sales['zipcode'], sales['price'])
plt.show()
```



98039 is the most expensive zip code.

8 Build a regression model with more features

```
In [21]: #my_features_model = (train_data, target='price', features=my_features, validation_set=None)
        x_train = train_data[my_features].values.reshape(-1, len(my_features))
        y_train = train_data['price'].values.reshape(-1, 1)
```

```
In [22]: mult_model = linear_model.LinearRegression()
        mult_model.fit(x_train, y_train)
```

```
Out[22]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

8.1 Comparing the results of the simple model with adding more features

```
In [23]: x_test = test_data[my_features].values.reshape(-1, len(my_features))
        y_test = test_data['price'].values.reshape(-1, 1)
```

```
y_pred = mult_model.predict(x_test)
```

```
In [24]: print('intercept:', mult_model.intercept_, 'coefficients:', mult_model.coef_)
        # The mean squared error
        print("RMSE: %.2f" % (rmse(y_pred, y_test)))
```

```
('intercept:', array([-56348418.9495597]), 'coefficients:', array([[ -6.17445455e+04,   1.880533
    -2.88006880e-01,  -9.68870610e+03,   5.75218487e+02]]))
```

```
RMSE: 249311.90
```

The RMSE goes down from \$254,323,39 to \$228,024,43 with more features.

9 Apply learned models to predict prices of 3 houses

The first house we will use is considered an "average" house in Seattle.

```
In [25]: house1 = sales[sales['id']==5309101200]
```

```
In [26]: house1
```

```
Out[26]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
1054	5309101200	20140605T000000	620000	4	2.25	2400	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
1054	5350	1.5	0	0	...	7	1460	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
1054	940	1929	0	98117	47.6763	-122.37	


```

      sqft_living15  sqft_lot15
1054           1250         4880

```

```
[1 rows x 21 columns]
```

```
In [27]: print house1['price']
```

```

1054    620000
Name: price, dtype: int64

```

```
In [28]: print simple_model.predict(house1['sqft_living'])
```

```
[[ 630959.4568039]]
```

```

/home/aluno/anaconda3/envs/gl-env/lib/python2.7/site-packages/sklearn/utils/validation.py:386: D
DeprecationWarning)

```

```
In [29]: print mult_model.predict(house1[my_features])
```

```
[[ 630924.33807747]]
```

In this case, the model with more features provides a worse prediction than the simpler model with only 1 feature. However, on average, the model with more features is better.

9.1 Prediction for a second, fancier house

We will now examine the predictions for a fancier house.

```
In [30]: house2 = sales[sales['id']==1925069082]
```

```
In [31]: house2
```

```

Out[31]:
      id      date  price  bedrooms  bathrooms  sqft_living  \
1361  1925069082  20150511T000000  2200000         5         4.25         4640

```

```

      sqft_lot  floors  waterfront  view  ...  grade  sqft_above  \
1361     22703     2.0           1     4  ...     8         2860

```

```

      sqft_basement  yr_built  yr_renovated  zipcode     lat     long  \
1361           1780     1952             0    98052  47.6393 -122.097

```

```

      sqft_living15  sqft_lot15
1361           3140         14200

```

```
[1 rows x 21 columns]
```

```
In [32]: print house2['price']
```

```
1361    2200000
```

```
Name: price, dtype: int64
```

```
In [33]: print simple_model.predict(house2['sqft_living'].values.reshape(-1,1))
```

```
[[ 1263248.45867172]]
```

```
In [34]: print mult_model.predict(house2[my_features])
```

```
[[ 1270172.16078085]]
```

In this case, the model with more features provides a better prediction. This behavior is expected here, because this house is more differentiated by features that go beyond its square feet of living space, especially the fact that it's a waterfront house.

9.2 Last house, super fancy

Our last house is a very large one owned by a famous Seattleite.

```
In [35]: bill_gates = {'bedrooms': [8],  
                      'bathrooms': [25],  
                      'sqft_living': [50000],  
                      'sqft_lot': [225000],  
                      'floors': [4],  
                      'zipcode': ['98039'],  
                      'condition': [10],  
                      'grade': [10],  
                      'waterfront': [1],  
                      'view': [4],  
                      'sqft_above': [37500],  
                      'sqft_basement': [12500],  
                      'yr_built': [1994],  
                      'yr_renovated': [2010],  
                      'lat': [47.627606],  
                      'long': [-122.242054],  
                      'sqft_living15': [5000],  
                      'sqft_lot15': [40000]}
```

```
In [36]: print simple_model.predict(pd.DataFrame(bill_gates)['sqft_living'].values.reshape(-1,1))
```

```
[[ 14067100.74649496]]
```

The model predicts a price of over \$14M for this house! But we expect the house to cost much more. (There are very few samples in the dataset of houses that are this fancy, so we don't expect the model to capture a perfect prediction here.)

```
In [37]: print mult_model.predict(pd.DataFrame(bill_gates)[my_features])

[[ 15779944.98847022]]
```

10 Answers

11 1 - Selection and summary statistics

```
In [38]: sales[sales['zipcode']==98039]
```

```
Out[38]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
2974	3625049014	20140829T000000	2950000	4	3.50	4860	
3761	2540700110	20150212T000000	1905000	4	3.50	4210	
4077	3262300940	20141107T000000	875000	3	1.00	1220	
4078	3262300940	20150210T000000	940000	3	1.00	1220	
4149	6447300265	20141014T000000	4000000	4	5.50	7080	
4411	2470100110	20140804T000000	5570000	5	5.75	9200	
4791	2210500019	20150324T000000	937500	3	1.00	1320	
5178	6447300345	20150406T000000	1160000	4	3.00	2680	
5589	6447300225	20141106T000000	1880000	3	2.75	2620	
5880	2525049148	20141007T000000	3418800	5	5.00	5450	
6868	3262300235	20141126T000000	1555000	5	2.50	2870	
7501	2525049133	20150402T000000	1398000	5	2.25	2640	
8241	3262301355	20140725T000000	1320000	3	2.75	2680	
9254	9208900037	20140919T000000	6885000	6	7.75	9890	
9694	3262301610	20141118T000000	865000	3	1.50	1530	
9809	5426300060	20141008T000000	1000000	3	2.25	2300	
11278	3025300226	20140515T000000	2100000	4	1.75	3550	
11952	2260300060	20150410T000000	2575000	5	3.00	4780	
12295	3738000070	20150309T000000	1712750	5	2.50	2660	
12811	5425700150	20140804T000000	787500	4	1.75	1580	
13235	3262300322	20150408T000000	1651000	4	3.25	3640	
13267	5425700205	20140520T000000	1800000	4	3.50	4460	
13419	2525049246	20141017T000000	1550000	2	2.25	2950	
13621	2525049266	20140821T000000	1762000	3	2.25	3060	
13988	5427110040	20140609T000000	1225000	4	2.50	2740	
14052	7397300220	20140529T000000	2750000	4	3.25	4430	
14254	2425049107	20150305T000000	1950000	4	3.75	4150	
14385	2425049061	20140825T000000	2200000	3	2.00	3570	
14803	3835502815	20140925T000000	1260000	3	2.50	3110	
15022	2210500010	20140930T000000	2450000	7	4.25	4670	
15255	2425049063	20140911T000000	3640900	4	3.25	4830	
15632	3625049088	20140702T000000	2271150	4	3.25	4040	
16268	3025300250	20150513T000000	1620000	4	2.25	2350	
16302	7397300170	20140530T000000	3710000	4	3.50	5550	
16377	3262300920	20150408T000000	1200000	4	3.00	2150	

16825	3025300095	20141009T000000	2500000	4	4.50	4300
17001	2525049259	20140812T000000	2187730	4	4.50	4240
17209	3025300225	20141031T000000	1450000	5	2.75	3090
17230	2470200020	20140514T000000	1880000	4	2.75	3260
17899	3262300555	20140708T000000	2458000	4	5.25	6500
17930	3625049079	20140801T000000	1350000	3	2.00	2070
18793	2525049263	20140709T000000	2680000	5	3.00	4290
18892	5427100150	20140626T000000	1410000	4	2.25	3250
18912	2425049066	20140616T000000	1920000	4	2.50	3070
19148	3625049042	20141011T000000	3635000	5	6.00	5490
19236	2525049086	20141003T000000	2720000	4	3.25	3990
19351	2525049113	20140725T000000	1950000	4	3.50	4065
20096	3262300485	20150421T000000	2250000	5	5.25	3410
21040	6447300365	20141113T000000	2900000	5	4.00	5190
21514	3262300818	20150227T000000	1865000	4	3.75	3790

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
2974	23885	2.0	0	0	...	12	4860	
3761	18564	2.0	0	0	...	11	4210	
4077	8119	1.0	0	0	...	7	1220	
4078	8119	1.0	0	0	...	7	1220	
4149	16573	2.0	0	0	...	12	5760	
4411	35069	2.0	0	0	...	13	6200	
4791	8500	1.0	0	0	...	7	1320	
5178	15438	2.0	0	2	...	8	2680	
5589	17919	1.0	0	1	...	9	2620	
5880	20412	2.0	0	0	...	11	5450	
6868	16238	2.0	0	0	...	8	2870	
7501	14959	1.0	0	0	...	7	1770	
8241	20104	1.0	0	0	...	9	1820	
9254	31374	2.0	0	4	...	13	8860	
9694	10827	1.0	0	0	...	8	1530	
9809	15952	1.0	0	0	...	8	1150	
11278	19865	2.0	0	0	...	9	3550	
11952	20440	1.0	0	0	...	10	3660	
12295	6572	1.0	0	0	...	9	1960	
12811	9382	1.0	0	0	...	7	1080	
13235	13530	1.0	0	0	...	9	2570	
13267	16953	1.0	0	0	...	9	2550	
13419	15593	1.0	0	0	...	8	1560	
13621	16000	2.0	0	0	...	10	3060	
13988	16007	2.0	0	0	...	9	2740	
14052	21000	2.0	0	0	...	10	4430	
14254	17424	1.0	0	0	...	9	3130	
14385	30456	1.0	0	1	...	8	2070	
14803	9930	1.0	0	1	...	8	1640	
15022	23115	2.0	0	2	...	11	4670	
15255	22257	2.0	1	4	...	11	4830	

15632	18916	1.0	0	0	...	9	4040
16268	17709	2.0	0	0	...	9	2350
16302	28078	2.0	0	2	...	12	3350
16377	8119	2.0	0	0	...	8	2150
16825	19844	2.0	0	0	...	11	4300
17001	13162	2.0	0	0	...	10	4240
17209	19865	1.0	0	0	...	9	3090
17230	19542	1.0	0	0	...	10	2170
17899	14986	2.0	0	0	...	11	5180
17930	9600	1.0	0	1	...	7	1590
18793	20445	2.0	0	0	...	11	4290
18892	16684	2.0	0	0	...	9	3250
18912	34412	1.0	0	3	...	9	2070
19148	19897	2.0	0	0	...	12	5490
19236	18115	2.0	0	0	...	11	3990
19351	18713	2.0	0	0	...	10	4065
20096	8118	2.0	0	0	...	11	3410
21040	14600	2.0	0	1	...	11	5190
21514	8797	2.0	0	0	...	11	3290

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long \
2974	0	1996	0	98039	47.6172	-122.230
3761	0	2001	0	98039	47.6206	-122.225
4077	0	1955	0	98039	47.6328	-122.236
4078	0	1955	0	98039	47.6328	-122.236
4149	1320	2008	0	98039	47.6151	-122.224
4411	3000	2001	0	98039	47.6289	-122.233
4791	0	1954	0	98039	47.6187	-122.226
5178	0	1902	1956	98039	47.6109	-122.226
5589	0	1949	0	98039	47.6144	-122.228
5880	0	2014	0	98039	47.6209	-122.237
6868	0	1962	0	98039	47.6308	-122.238
7501	870	1929	0	98039	47.6191	-122.234
8241	860	1964	0	98039	47.6304	-122.234
9254	1030	2001	0	98039	47.6305	-122.240
9694	0	1955	0	98039	47.6354	-122.234
9809	1150	1963	0	98039	47.6322	-122.232
11278	0	1962	2002	98039	47.6236	-122.235
11952	1120	1975	0	98039	47.6242	-122.239
12295	700	1959	0	98039	47.6176	-122.223
12811	500	1963	0	98039	47.6353	-122.232
13235	1070	1924	2000	98039	47.6293	-122.238
13267	1910	1962	1994	98039	47.6338	-122.232
13419	1390	1942	1986	98039	47.6209	-122.236
13621	0	1988	0	98039	47.6189	-122.230
13988	0	1984	0	98039	47.6353	-122.229
14052	0	1952	2007	98039	47.6398	-122.237
14254	1020	1963	2000	98039	47.6390	-122.236

14385	1500	1946	1982	98039	47.6413	-122.240
14803	1470	1954	0	98039	47.6112	-122.226
15022	0	1992	0	98039	47.6183	-122.227
15255	0	1990	0	98039	47.6409	-122.241
15632	0	1954	0	98039	47.6155	-122.238
16268	0	1977	0	98039	47.6232	-122.236
16302	2200	2000	0	98039	47.6395	-122.234
16377	0	1953	2004	98039	47.6335	-122.236
16825	0	1985	1999	98039	47.6218	-122.237
17001	0	2004	0	98039	47.6193	-122.229
17209	0	1953	0	98039	47.6232	-122.235
17230	1090	1968	0	98039	47.6245	-122.236
17899	1320	2001	0	98039	47.6304	-122.236
17930	480	1946	0	98039	47.6160	-122.239
18793	0	1985	0	98039	47.6217	-122.239
18892	0	1979	0	98039	47.6334	-122.229
18912	1000	1950	0	98039	47.6400	-122.240
19148	0	2005	0	98039	47.6165	-122.236
19236	0	1989	0	98039	47.6177	-122.229
19351	0	1987	0	98039	47.6209	-122.237
20096	0	2006	0	98039	47.6295	-122.236
21040	0	2013	0	98039	47.6102	-122.225
21514	500	2006	0	98039	47.6351	-122.236

	sqft_living15	sqft_lot15
2974	3580	16054
3761	3520	18564
4077	1910	8119
4078	1910	8119
4149	3140	15996
4411	3560	24345
4791	2790	10800
5178	4480	14406
5589	3400	14400
5880	3160	17825
6868	2870	16238
7501	3240	17904
8241	3060	19837
9254	4540	42730
9694	2050	10827
9809	2200	14284
11278	3000	19862
11952	4660	20440
12295	3960	14595
12811	2010	9382
13235	2760	15000
13267	1980	13370
13419	2060	19855

13621	3510	13162
13988	2760	16008
14052	3930	20000
14254	3930	21420
14385	3570	27418
14803	3650	14399
15022	3240	13912
15255	3820	25582
15632	3000	18831
16268	3360	19855
16302	2980	19602
16377	1590	8119
16825	3070	19845
17001	3010	12163
17209	2970	19862
17230	3480	19863
17899	2270	8119
17930	3000	16215
18793	3620	22325
18892	2890	16927
18912	3780	27940
19148	2910	17600
19236	3450	16087
19351	3070	18713
20096	3410	16236
21040	3840	19250
21514	2660	12150

[50 rows x 21 columns]

```
In [39]: sales[sales['zipcode']==98039]['price'].mean()
```

```
Out[39]: 2160606.6
```

12 2 - Filtering data

```
In [40]: num = len(sales[(sales['sqft_living']>2000) & (sales['sqft_living']<4000)])
num
```

```
Out[40]: 9111
```

```
In [41]: num/float(len(sales))
```

```
Out[41]: 0.4215518437977143
```

13 3 - Building a regression model with several more features

```
In [42]: advanced_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode',
                             'waterfront', 'view', 'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovate']
```

```

        'sqft_living15', 'sqft_lot15']
x_train2 = train_data[advanced_features].values.reshape(-1, len(advanced_features))
y_train2 = train_data['price'].values.reshape(-1, 1)

In [43]: mult_model = linear_model.LinearRegression()
         mult_model.fit(x_train2, y_train2)

Out[43]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [44]: x_test2 = test_data[advanced_features].values.reshape(-1, len(advanced_features))
         y_test2 = test_data['price'].values.reshape(-1, 1)

         y_pred2 = mult_model.predict(x_test2)

In [45]: print('intercept:', simple_model.intercept_, 'coefficients:', simple_model.coef_)
         # The mean squared error
         print("RMSE: %.2f" % (rmse(y_pred2, y_test2)))

('intercept:', array([-46493.04519733]), 'coefficients:', array([[ 282.27187583]]))
RMSE: 193713.58

In [46]: print('intercept:', mult_model.intercept_, 'coefficients:', mult_model.coef_)
         # The mean squared error
         print("RMSE: %.2f" % (rmse(y_pred, y_test)))

('intercept:', array([ 7220465.95346491]), 'coefficients:', array([[ -3.85547540e+04,   4.429332
 1.59043347e-01,   4.14361852e+03,  -6.01513553e+02,
 2.44382135e+04,   9.34950344e+04,   6.12210257e+05,
 5.06076502e+04,   6.98882157e+01,   4.40768912e+01,
-2.62156390e+03,   1.96225739e+01,   6.03747455e+05,
-2.25535679e+05,   2.24949849e+01,  -4.25865615e-01]]))
RMSE: 249311.90

In [52]: print('Difference: ', (rmse(y_pred, y_test)) - (rmse(y_pred2, y_test2)))

('Difference: ', 55598.322088692774)

```